

# Docker Networking

# Agenda



## In this session, you will learn about:

- Networking Overview
- Understanding CNM
- Network Drivers

# Scope of this section

- This section does not go into **OS-specific** details about how Docker networking works.
- Example:
  - Linux Bridge
  - iptables
  - Packet Encapsulation/Decapsulation



# Networking overview

- Docker containers uses **Linux Bridge** feature by default to communicate.
- Whether your Docker hosts run Linux, Windows, or a mix of the two, you can use Docker to manage them in a platform-agnostic way.

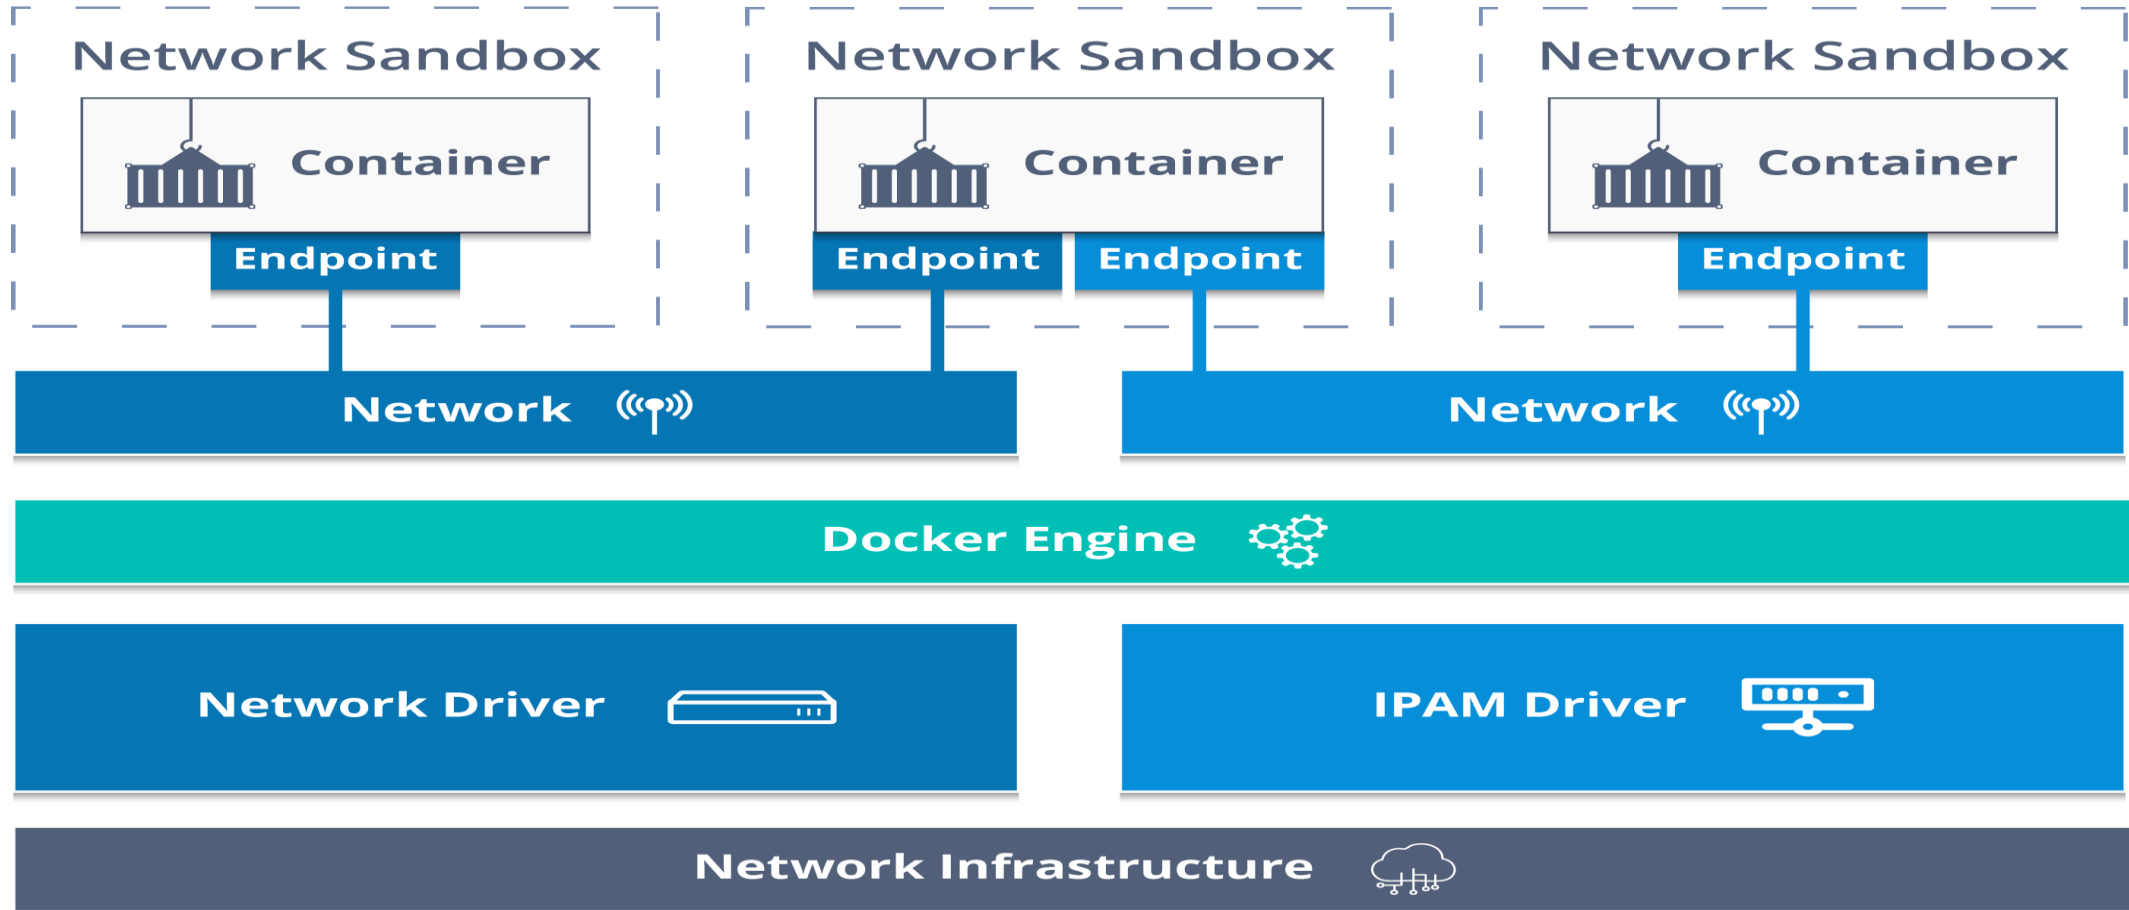
# CNM – Container Network Model

- **CNM** is an open-source networking specification for containers.
- **CNM** defines Networks, Endpoints, Sandboxes.
- **Libnetwork** is Docker's implementation of CNM.
- **Libnetwork** is extensible via pluggable drivers which allows various network technologies.

# Containers and CNM

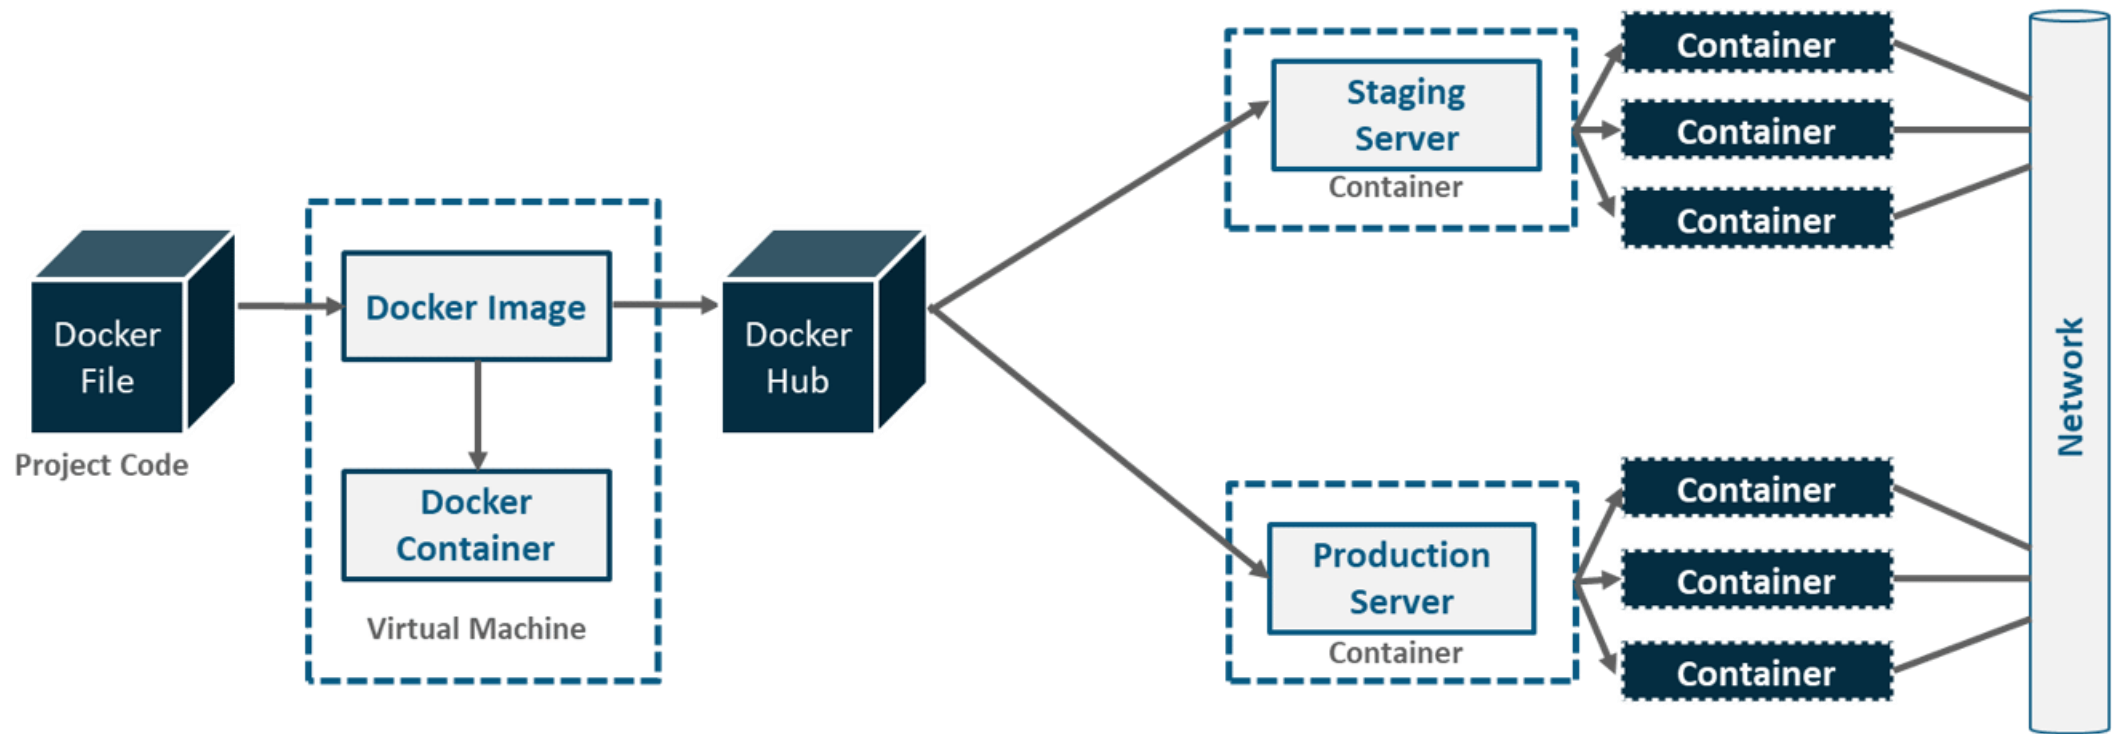


# Architecture of Container Network Model



Architecture of Container Networking Model

# Networking - Workflow





# CNM built on 5 Objects

- A CNM has mainly built on 5 objects:
  - Network Controller
  - Driver
  - Network
  - Endpoint
  - Sandbox

# Network Controller

Provides the entry-point into **Libnetwork** that exposes simple APIs for Docker Engine to allocate and manage networks.

# Driver

- Owns the Network
- Responsible for managing the Network
- Supports multiple drivers to satisfy various use-cases and deployment scenarios.

# Network

- Provides connectivity between a group of endpoints that belong to the same network and isolate from the rest.
- Whenever a network is created or updated, the corresponding Driver will be notified of the event.

# Endpoint

- Provides the connectivity for container in a Network.

# Sandbox

- Created when users request to create an endpoint on a network.
- A Sandbox can have multiple endpoints attached to different networks representing container's network configuration such as:
  - IP-Address
  - MAC-Address
  - Routes
  - DNS

# Network Drivers

# Network Drivers

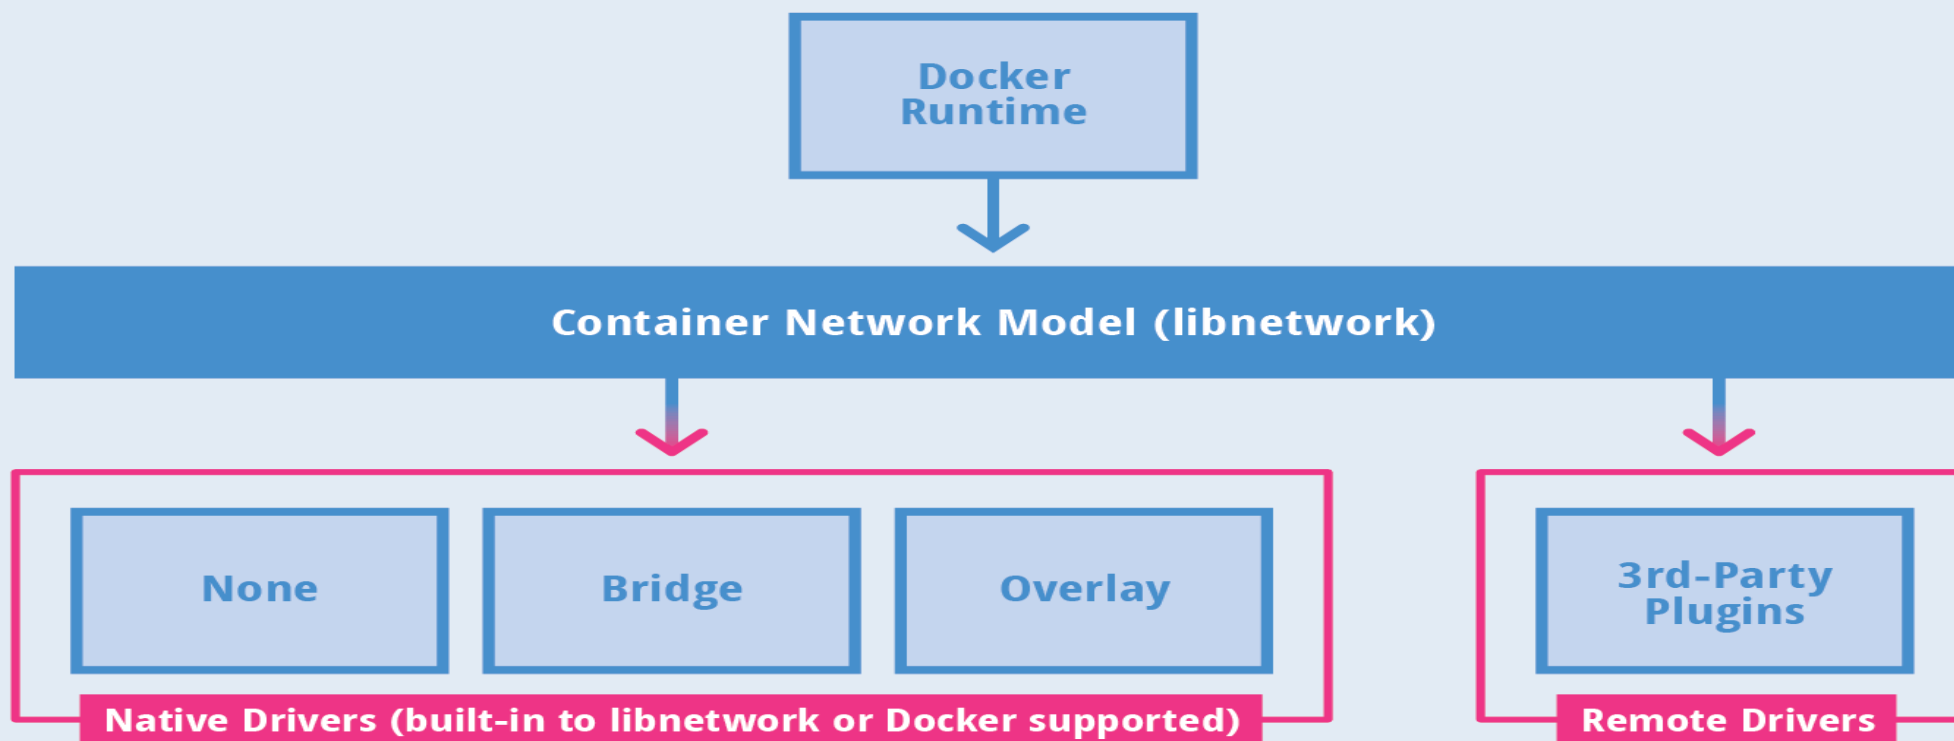
- bridge
- host
- overlay
- macvlan
- none

```
$ docker info  
  
<snip>  
...  
Plugins:  
Volume: local  
Network: bridge host macvlan none overlay  
...  
<snip>
```



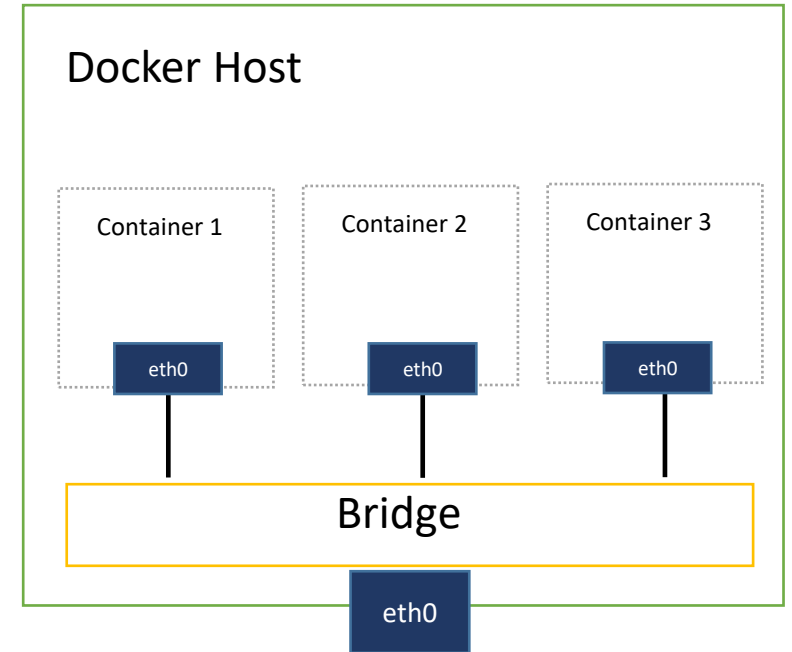
# Network Drivers

## Container Network Model (CNM) Drivers



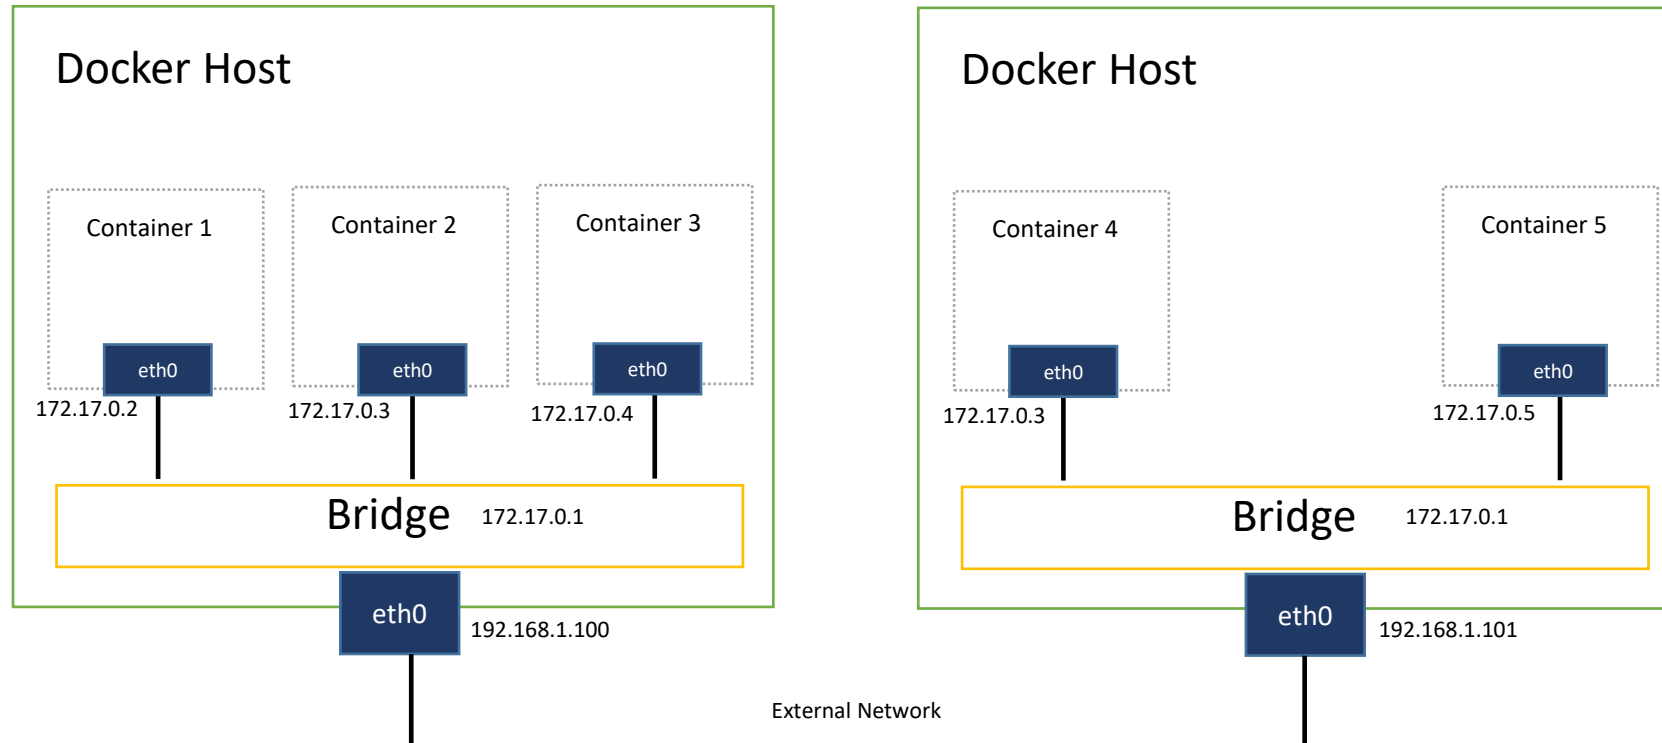
# Bridge Driver

- Default driver when a container is run
- Additional bridge networks can be created
- Bridge networks apply to containers running on the same host
- Has its own internal IP address allocation mechanism
- Example : 172.17.0.0/16 (Default)
- Very robust



# Bridge Driver

- Bridge networks are usually used when your applications run in standalone containers that need to communicate.
- Container to external network done via PAT (iptables)



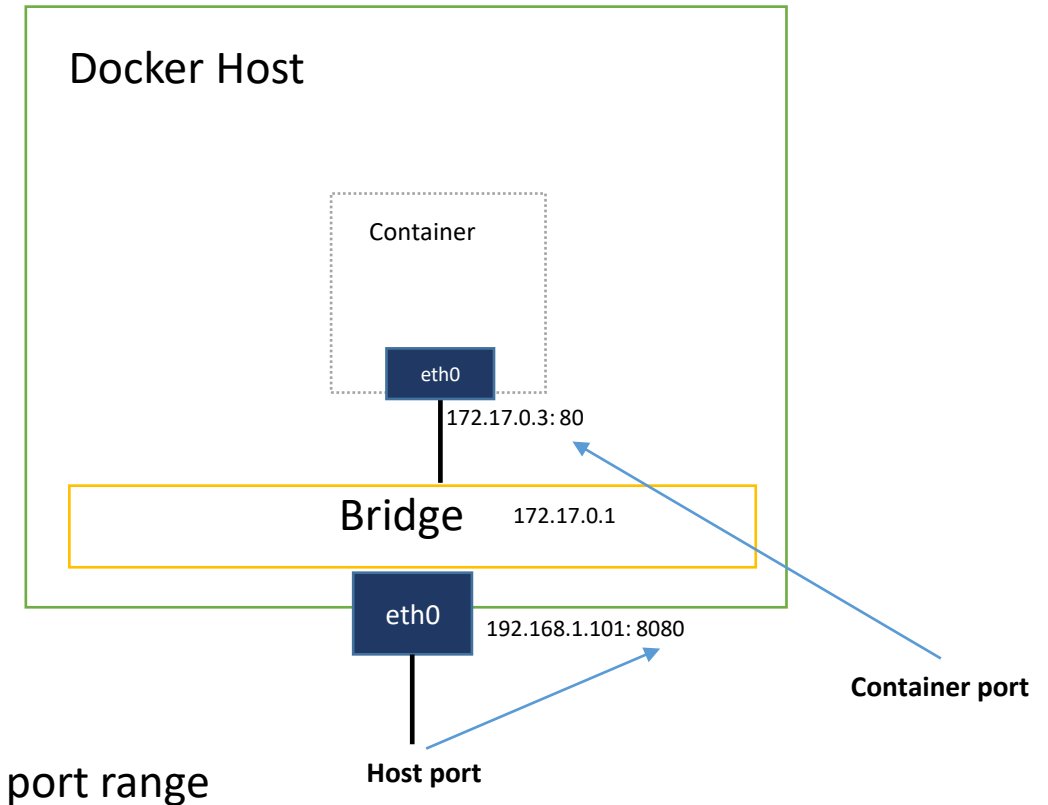
# Bridge Networking and Port Mapping

- Access to containers via port mapping
- Map container port to host port

**`docker container run -p 8080:80 nginx`**

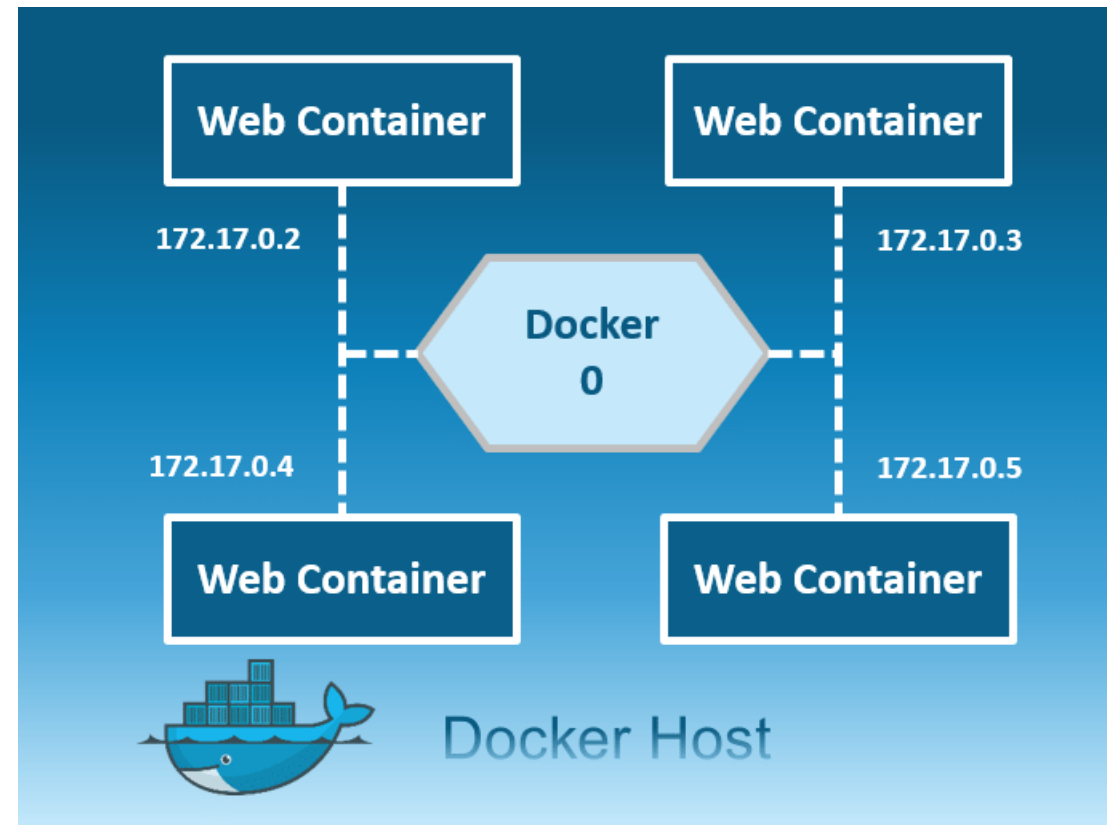
Host Port – 8080

Container Port - 80



```
# cat /proc/sys/net/ipv4/ip_local_port_range  
32768 to 60999
```

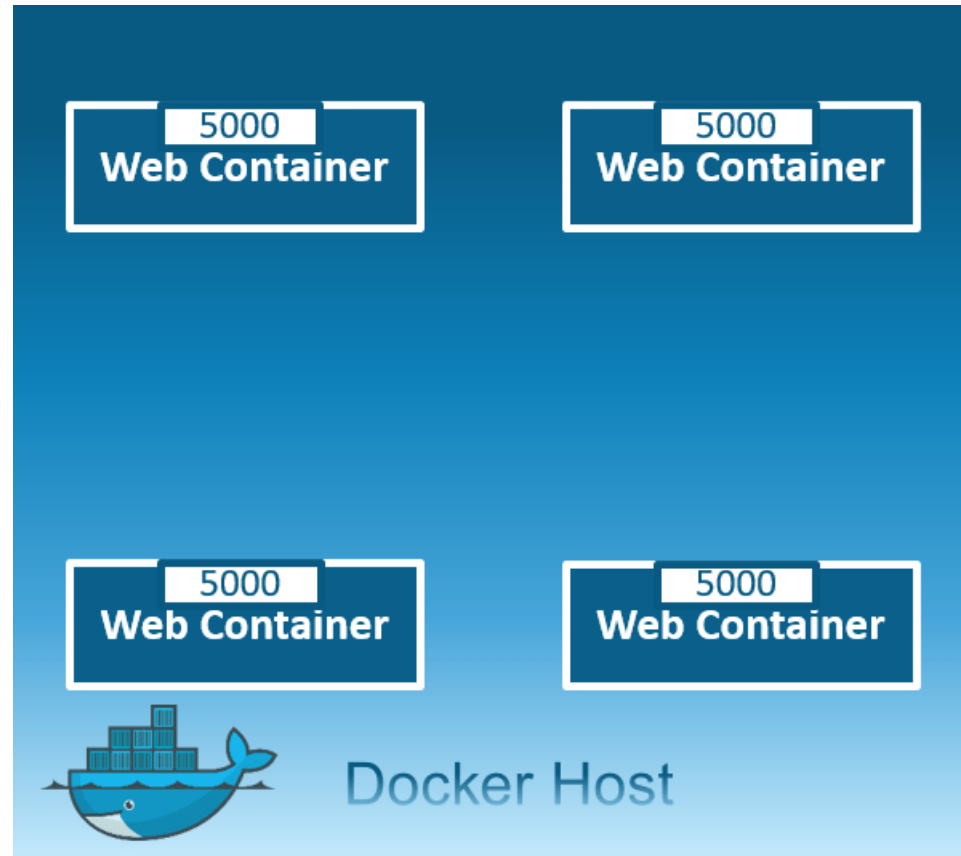
# Example – Bridge driver



# Host

- Driver removes the network isolation between the docker host and the docker containers to use the host's networking directly.
- Will not be able to run multiple web containers on the same host, on the same port as the port is now common to all containers in the host network.

# Example – Host driver

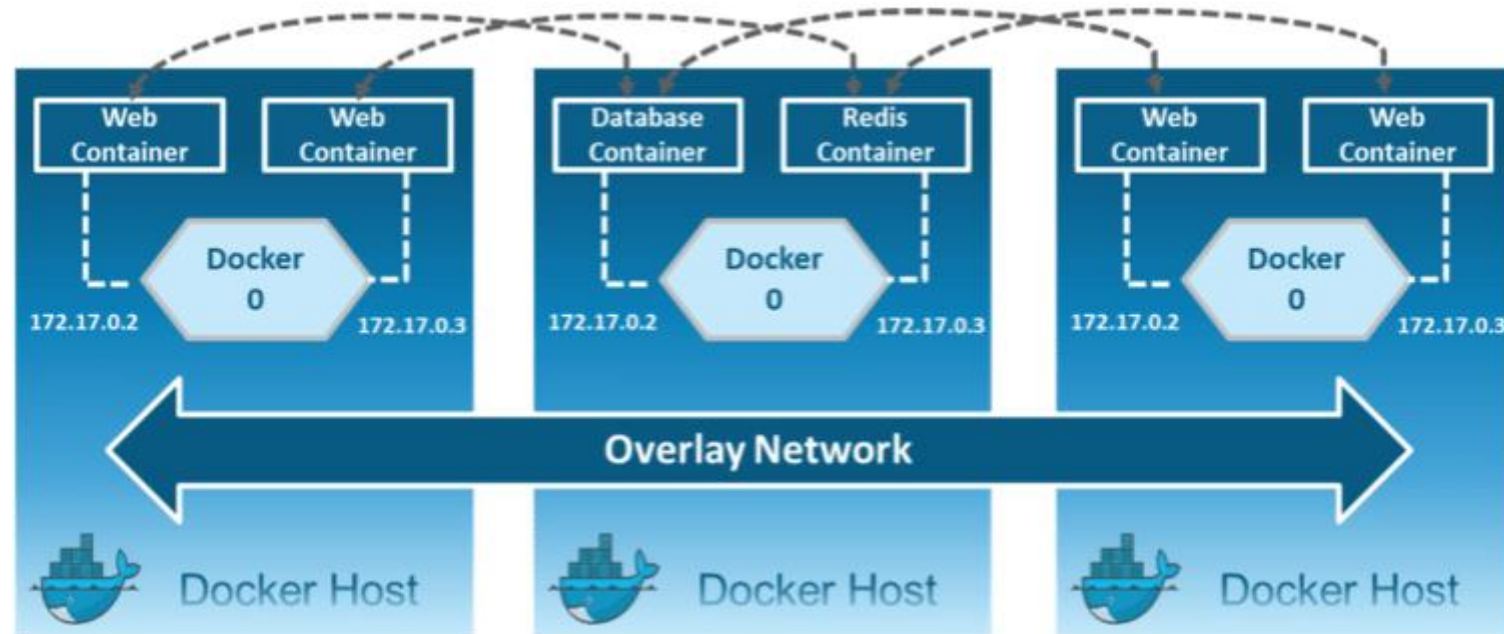


# Overlay

- Creates an internal private network that spans across all the nodes participating in the swarm cluster.
- So, Overlay networks facilitate communication between a swarm service and a standalone container, or between two standalone containers on different Docker Daemons.



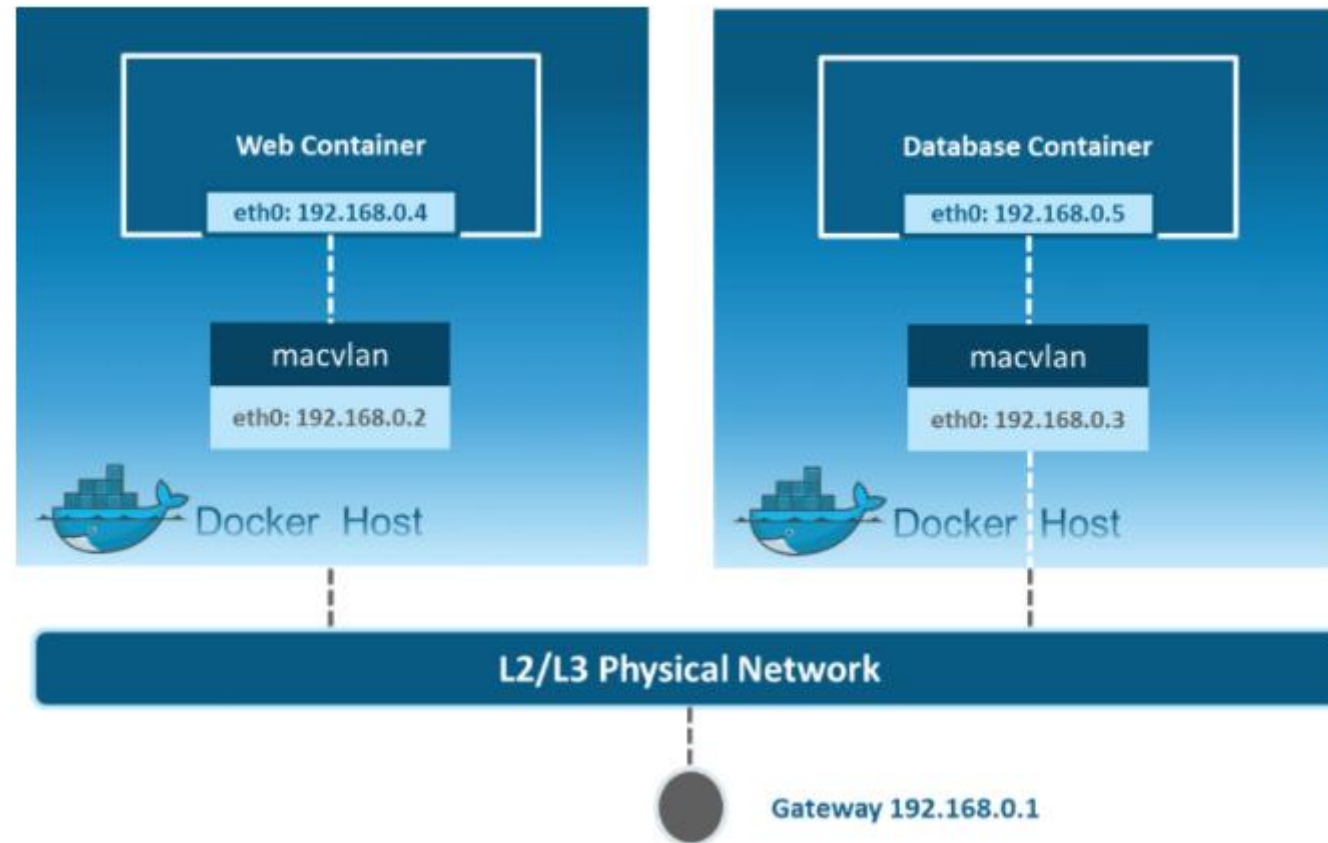
# Example – Overlay driver



# Macvlan

- Allows you to assign a MAC address to a container, making it appear as a physical device on your network.
- Docker daemon routes traffic to containers by their MAC addresses.
- Macvlan driver is the best choice when you are expected to be directly connected to the physical network, rather than routed through the Docker host's network stack.

# Example – Macvlan driver



# Disable Networking for Container

- Use the **--network none** flag when starting the container to disable networking
- Within the container, only the loopback device is created



# Networking Drivers Summary

- **Bridge networks** - Best when you need multiple containers to communicate on the same Docker host.
- **host**: Uses the host's networking directly.
- **none**: Disables all networking.
- **Overlay networks** - Best when you need containers running on different Docker hosts to communicate, or when multiple applications work together using swarm services.
- **Macvlan networks** - Best when you are migrating from a VM setup or need your containers to look like physical hosts on your network, each with a unique MAC address.
- **Third-party network plugins** - allow you to integrate Docker with specialized network stacks.

# Basic Docker Network Commands

- `docker network ls`
- `docker network create`
- `docker network inspect`
- `docker network rm`

# Lab – Docker Networking