

# Lab: Working with Deployments

## Introduction:

Deployments represent a set of multiple, identical Pods with no unique identities. Deployment runs multiple replicas of your application and automatically replaces any instances that fail or become unresponsive.

In this way, Deployments help ensure that one or more instances of your application are available to serve user requests. Deployments are managed by the Kubernetes Deployment controller.

Deployment ensures that only a certain number of Pods are down while they are being updated. By default, it ensures that at least **75%** of the desired number of Pods are up (25% max unavailable).

## Objectives:

- Create a Deployment with Imperative Method
- Generate the Deployment manifest in the yml & json format
- Create a Deployment with Declarative Method
- Exposing Deployment to a Service i.e., ClusterIP
- Scale-Up/Down Deployment
- Cleanup

Ensure that you have logged-in as **root** user on **eoc-controller** node.

## 1. Create a Deployment with Imperative Method

**1.1** Let's create deployment named **redis-deployment** and using the Redis image.

```
# kubectl create --help
```

```
# kubectl create deployment --help
```

```
# kubectl create deployment redis-deployment --image=redis
```

## Output:

```
[root@eoc-controller ~]# kubectl create deployment redis-deployment --image=redis
deployment.apps/redis-deployment created
```

1.2 Let's check the details of the deployment by executing the command.

```
# kubectl describe deployment redis-deployment
```

Output:

```
[root@eoc-controller ~]#kubectl describe deployment redis-deployment
Name:                redis-deployment
Namespace:            default
CreationTimestamp:    Tue, 05 Sep 2023 09:49:06 -0400
Labels:               app=redis-deployment
Annotations:          deployment.kubernetes.io/revision: 1
Selector:              app=redis-deployment
Replicas:              1 desired | 1 updated | 1 total | 1 available | 0 unavailable
StrategyType:          RollingUpdate
MinReadySeconds:       0
RollingUpdateStrategy: 25% max unavailable, 25% max surge
Pod Template:
  Labels:  app=redis-deployment
  Containers:
    redis:
      Image:        redis
      Port:          <none>
      Host Port:     <none>
      Environment:   <none>
      Mounts:         <none>
      Volumes:        <none>
  Conditions:
    Type           Status  Reason
    ----           -
    Available       True    MinimumReplicasAvailable
    Progressing     True    NewReplicaSetAvailable
OldReplicaSets: <none>
```

1.3 Let's list the pods by executing the below command.

```
# kubectl get pods
```

Output:

```
[root@eoc-controller ~]#kubectl get pods
NAME                                READY   STATUS    RESTARTS   AGE
redis-deployment-785c545fd6-znv97  1/1     Running   0           5m36s
```

1.4 Let's delete the deployment by executing the below command.

```
# kubectl delete deployment redis-deployment
```

Output:

```
[root@eoc-controller ~]#kubectl delete deployment redis-deployment
deployment.apps "redis-deployment" deleted
```

1.5 Let's create another deployment of nginx with 4 replicas.

```
# kubectl create deployment nginx-deployment --image=nginx \
--replicas=4 --dry-run=client -o yaml | tee nginx-
deployment.yaml
```

### Output:

```
[root@eoc-controller ~]# kubectl create deployment nginx-deployment --image=nginx \
> --replicas=4 --dry-run=client -o yaml | tee nginxdeployment.yml
apiVersion: apps/v1
kind: Deployment
metadata:
  creationTimestamp: null
  labels:
    app: nginx-deployment
  name: nginx-deployment
spec:
  replicas: 4
  selector:
    matchLabels:
      app: nginx-deployment
  strategy: {}
  template:
    metadata:
      creationTimestamp: null
      labels:
        app: nginx-deployment
    spec:
      containers:
      - image: nginx
        name: nginx
        resources: {}
status: {}
```

1.6 Let's use the **nginx-deployment.yml** file created from the above step for creating deployment.

```
# kubectl apply -f nginx-deployment.yml
```

### Output:

```
[root@eoc-controller ~]# kubectl apply -f nginx-deployment.yml
deployment.apps/nginx-deployment created
```

1.7 Let's list the pods by executing the below command.

```
# kubectl get pods
```

### Output:

```
[root@eoc-controller ~]# kubectl get pods
```

NAME	READY	STATUS	RESTARTS	AGE
nginx-deployment-6d6565499c-dmjcf	1/1	Running	0	35s
nginx-deployment-6d6565499c-h9z15	1/1	Running	0	35s
nginx-deployment-6d6565499c-l5tvp	1/1	Running	0	35s
nginx-deployment-6d6565499c-r2b44	1/1	Running	0	35s

1.8 Let's **delete** the deployment and notice that it deletes the pods as well.

```
# kubectl delete -f nginx-deployment.yml
```

Output:

```
[root@eoc-controller ~]#kubectl delete -f nginx-deployment.yml
deployment.apps "nginx-deployment" deleted
```

## 2. Create a Deployment with Declarative Method

2.1 Let's view the yaml manifest file.

```
# cat -n ~/kubernetes/deployment-webserver.yml
```

Output:

```
[root@eoc-controller ~]#cat -n ~/kubernetes/deployment-webserver.yml
 1  apiVersion: apps/v1
 2  kind: Deployment
 3  metadata:
 4    name: web-server
 5  spec:
 6    replicas: 4
 7    selector:
 8      matchLabels:
 9        tier: web-server
10  template:
11    metadata:
12      labels:
13        tier: web-server
14    spec:
15      containers:
16      - name: web-server-container
17        image: nginx:1.19
```

2.2 Let's create the deployment using the deployment-webserver.yml file.

```
# kubectl apply -f ~/kubernetes/deployment-webserver.yml
```

Output:

```
[root@eoc-controller ~]#kubectl apply -f ~/kubernetes/deployment-webserver.yml
deployment.apps/web-server created
```

2.3 Let's list the deployment by executing the below command.

```
# kubectl get deployment
```

Output:

```
[root@eoc-controller ~]#kubectl get deployment
NAME          READY    UP-TO-DATE    AVAILABLE    AGE
web-server    4/4      4             4            25s
```

**2.4** Let's list the **labels** automatically generated for each Pod by executing the below command.

```
# kubectl get pods --show-labels
```

**Output:**

```
[root@eoc-controller ~]# kubectl get pods --show-labels
NAME                                READY    STATUS    RESTARTS   AGE    LABELS
web-server-5dcf74945f-2xnkq        1/1      Running   0           70s    pod-template-hash=5dcf74945f,tier=web-server
web-server-5dcf74945f-7xkrf        1/1      Running   0           70s    pod-template-hash=5dcf74945f,tier=web-server
web-server-5dcf74945f-f5c6g        1/1      Running   0           70s    pod-template-hash=5dcf74945f,tier=web-server
web-server-5dcf74945f-vq2gp        1/1      Running   0           70s    pod-template-hash=5dcf74945f,tier=web-server
```

**2.5** Let's list the ReplicaSet by executing the below command.

```
# kubectl get rs
```

**Output:**

```
[root@eoc-controller ~]# kubectl get rs
NAME                                DESIRED    CURRENT    READY    AGE
web-server-5dcf74945f              4           4           4        119s
```

**3.** Let's **expose** the deployment to the service to access the application.

```
# kubectl expose deployment web-server --name \
demo-service --type ClusterIP --labels app=nginx --port 80
```

**Output:**

```
[root@eoc-controller ~]# kubectl expose deployment web-server --name \
> demo-service --type ClusterIP --labels app=nginx --port 80
service/demo-service exposed
```

**3.1** Let's list the service by executing the below command.

```
# kubectl get svc demo-service
```

**Output:**

```
[root@eoc-controller ~]# kubectl get svc demo-service
NAME            TYPE        CLUSTER-IP    EXTERNAL-IP    PORT(S)    AGE
demo-service    ClusterIP   10.96.124.168 <none>         80/TCP     28s
```

**3.2** Let's **access** the application using the cluster-ip by executing the below command.

```
# CIP=`kubectl get svc demo-service -o
jsonpath={.spec.clusterIP}`
```

```
# kubectl get svc demo-service -o jsonpath={.spec.clusterIP}
```

```
# curl $CIP
```

```
# echo $CIP
```

Output:

```
[root@eoc-controller ~]#echo $CIP
10.96.124.168
```

Output:

```
[root@eoc-controller ~]#curl $CIP
<!DOCTYPE html>
<html>
<head>
<title>Welcome to nginx!</title>
<style>
  body {
    width: 35em;
    margin: 0 auto;
    font-family: Tahoma, Verdana, Arial, sans-serif;
  }
</style>
</head>
<body>
<h1>Welcome to nginx!</h1>
<p>If you see this page, the nginx web server is successfully installed and
working. Further configuration is required.</p>

<p>For online documentation and support please refer to
<a href="http://nginx.org/">nginx.org</a>.<br/>
Commercial support is available at
<a href="http://nginx.com/">nginx.com</a>.</p>

<p><em>Thank you for using nginx.</em></p>
</body>
</html>
```

4. Let's **scale up** the deployment by executing the below command.

```
# kubectl scale deployment web-server --replicas=6
```

Output:

```
[root@eoc-controller ~]#kubectl scale deployment web-server --replicas=6
deployment.apps/web-server scaled
```

4.1 Let's **verify** the self-healing property assured by the deployment.

```
# kubectl delete pod --all
```

Output:

```
[root@eoc-controller ~]#kubectl delete pod --all
pod "web-server-5dcf74945f-2xnkq" deleted
pod "web-server-5dcf74945f-54zwt" deleted
pod "web-server-5dcf74945f-7xkrf" deleted
pod "web-server-5dcf74945f-f5c6g" deleted
pod "web-server-5dcf74945f-t79fx" deleted
pod "web-server-5dcf74945f-vq2gp" deleted
```

4.2 Let's list the ReplicaSet by executing the below command.

```
# kubectl get rs
```

Output:

```
[root@eoc-controller ~]#kubectl get rs
NAME                                DESIRED    CURRENT    READY    AGE
web-server-5dcf74945f              6          6          6        6m10s
```

4.3 Lets scale down to 2 replicas by executing below command.

```
# kubectl scale deployment web-server --replicas=2
```

Output:

```
[root@eoc-kubemaster ~]#kubectl scale deployment web-server --replicas=2
deployment.apps/web-server scaled
```

4.4 Let's list the ReplicaSet by executing the below command.

```
# kubectl get rs
```

Output:

```
[root@eoc-controller ~]#kubectl scale deployment web-server --replicas=2
deployment.apps/web-server scaled
```

## 5 Cleanup

5.1 Let's delete the deployment by executing below command.

```
# kubectl delete deployment web-server
```

Output:

```
[root@eoc-controller ~]#kubectl delete deployment web-server
deployment.apps "web-server" deleted
```

5.2 Let's delete the service by executing below command.

```
# kubectl delete svc demo-service
```

Output:

```
[root@eoc-controller ~]#kubectl delete svc demo-service
service "demo-service" deleted
```