

Lab: Horizontal Pod Autoscaler

Introduction:

The Horizontal Pod Autoscaler automatically scales the number of pods in a replication controller, deployment, replica set or stateful set based on Observed CPU utilization. The Horizontal Pod Autoscaler is implemented as a Kubernetes API resource and a controller.

The resource determines the behaviour of the controller. The controller periodically adjusts the number of replicas in a replication controller or deployment to match the observed average CPU utilization to the target specified by user.

Objectives:

- Create A Nginx Deployment
- Create HorizontalPodAutoscalers
- Install siege -http Load Simulator
- Watch Auto Scale Up & Watch Auto Scale Down
- Cleanup

Ensure that you have logged-in as **root** user on **eoc-controller** node.

1. Create A Nginx Deployment

1.1 Let's **view** the manifest to create a deployment for nginx by executing the below command.

```
# cat -n ~/kubernetes/hpa-deployment.yml
```

Output:

```
[root@eoc-controller ~]#cat -n ~/kubernetes/hpa-deployment.yml
 1  apiVersion: apps/v1
 2  kind: Deployment
 3  metadata:
 4    name: web-application
 5  spec:
 6    replicas: 1
 7    selector:
 8      matchLabels:
 9        app: web
10  template:
11    metadata:
12      labels:
13        app: web
14    spec:
15      containers:
16      - image: nginx
17        name: web-application-container
18      resources:
19        limits:
20          cpu: "100m"
21        requests:
22          cpu: "100m"
```

Note: The cpu resource is limited to 100m this will help us to scale based on cpu utilization.

1.2 Let's create the deployment by executing the below command.

```
# kubectl create -f ~/kubernetes/hpa-deployment.yml
```

Output:

```
[root@eoc-controller ~]#kubectl apply -f ~/kubernetes/hpa-deployment.yml
deployment.apps/web-application created
```

1.3 Let's verify the deployment by executing the below command.

```
# kubectl get all -l app=web
```

Output:

```
[root@eoc-controller ~]#kubectl get all -l app=web
NAME                                READY    STATUS    RESTARTS    AGE
pod/web-application-6f669d8fb6-bf8n2  1/1      Running   0            29s

NAME                                DESIRED    CURRENT    READY    AGE
replicaset.apps/web-application-6f669d8fb6  1          1          1        29s
```

1.4 Let's **create** and expose a service of ClusterIP type to access our pod.

```
# kubectl expose deployment web-application --name \
web-service --port=80
```

Output:

```
[root@eoc-controller ~]# kubectl expose deployment web-application --name \
> web-service --port=80
service/web-service exposed
```

1.5 Let's **verify** the service by executing the below command.

```
# kubectl get svc web-service
```

Output:

```
[root@eoc-controller ~]# kubectl get svc web-service
NAME                TYPE          CLUSTER-IP    EXTERNAL-IP  PORT(S)    AGE
web-service         ClusterIP     10.105.203.79 <none>       80/TCP     2m6s
```

2 Create HorizontalPodAutoscalers

2.1 Let's **view** the manifest for creating hpa by executing the below command.

```
# cat -n ~/kubernetes/hpa-configure.yml
```

Output:

```
[root@eoc-controller ~]# cat -n ~/kubernetes/hpa-configure.yml
1  apiVersion: autoscaling/v1
2  kind: HorizontalPodAutoscaler
3  metadata:
4    name: demo-hpa
5  spec:
6    maxReplicas: 5
7    minReplicas: 1
8    scaleTargetRef:
9      apiVersion: apps/v1
10     kind: Deployment
11     name: web-application
12     targetCPUUtilizationPercentage: 20
```

2.2 Let's **create** the hpa by executing the below command.

```
# kubectl apply -f ~/kubernetes/hpa-configure.yml
```

Output:

```
[root@eoc-controller ~]# kubectl apply -f ~/kubernetes/hpa-configure.yml
horizontalpodautoscaler.autoscaling/demo-hpa created
```

2.3 Let's **verify** hpa details by executing the below command.

```
# kubectl get hpa
```

Output:

```
[root@eoc-controller ~]# kubectl get hpa
NAME          REFERENCE                TARGETS  MINPODS  MAXPODS  REPLICAS  AGE
demo-hpa      Deployment/web-application  0%/20%   1         5         1          94s
```

Info: Lets now simulate the load to increase the cpu utilization and watch the hpa work. We will use **siege** – which is an powerful HTTP load testing and benchmarking utility.

3 Install siege -http Load Simulator

3.1 Let's **install** the **epel repository** by executing below commands.

```
# dnf install -y epel-release
```

Output:

```
[root@eoc-controller ~]# dnf install -y epel-release
Last metadata expiration check: 19:56:04 ago on Thu 07 Sep 2023 07:09:44 AM EDT.
Dependencies resolved.
=====
Package                Architecture    Version           Repository        Size
=====
Installing:
epel-release            noarch          8-11.el8          extras            24 k
Transaction Summary
=====
Install 1 Package

Total download size: 24 k
Installed size: 35 k
Downloading Packages:
epel-release-8-11.el8.noarch.rpm              76 kB/s | 24 kB    00:00
-----
Total                                          75 kB/s | 24 kB    00:00
Running transaction check
Transaction check succeeded.
Running transaction test
Transaction test succeeded.
```

3.2 Let's **install** the **siege** utility by executing below commands.

```
# dnf install -y siege.x86_64
```

Output:

```
[root@eoc-controller ~]#dnf install -y siege.x86_64
Last metadata expiration check: 0:01:19 ago on Fri 08 Sep 2023 03:06:08 AM EDT.
Dependencies resolved.
=====
Package                Architecture      Version           Repository        Size
=====
Installing:
siege                  x86_64            4.1.2-1.el8       epel               121 k
Installing dependencies:
libjoedog              x86_64            0.1.2-13.el8      epel               25 k
Transaction Summary
=====
Install 2 Packages

Total download size: 146 k
Installed size: 320 k
Downloading Packages:
(1/2): libjoedog-0.1.2-13.el8.x86_64.rpm      199 kB/s | 25 kB    00:00
(2/2): siege-4.1.2-1.el8.x86_64.rpm           773 kB/s | 121 kB   00:00
-----
Total                                           128 kB/s | 146 kB   00:01
Extra Packages for Enterprise Linux 8 - x86_64 1.6 MB/s | 1.6 kB   00:00
Importing GPG key 0x2F86D6A1:
```

3.3 Let's verify that **siege** is installed correctly by executing the below command.

```
# siege -V
```

Output:

```
[root@eoc-controller ~]#siege -V
New configuration template added to /root/.siege
Run siege -C to view the current settings in that file
SIEGE 4.1.2

Copyright (C) 2022 by Jeffrey Fulmer, et al.
This is free software; see the source for copying conditions.
There is NO warranty; not even for MERCHANTABILITY or FITNESS
FOR A PARTICULAR PURPOSE.
```

Note: Open another terminal and run the below command.

4 Watch Auto Scale Up & Watch Auto Scale Down

4.1 Let's watch the deployment by executing the below command.

Terminal 2:

```
# watch kubectl get all
```

Output:

```
Every 2.0s: kubectl get all                                eoc-controller: Fri Sep  8 03:09:56 2023
```

NAME	READY	STATUS	RESTARTS	AGE
pod/web-application-6f669d8fb6-bf8n2	1/1	Running	0	14m

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
service/kubernetes	ClusterIP	10.96.0.1	<none>	443/TCP	4d
service/web-service	ClusterIP	10.105.203.79	<none>	80/TCP	13m

NAME	READY	UP-TO-DATE	AVAILABLE	AGE
deployment.apps/web-application	1/1	1	1	14m

NAME	DESIRED	CURRENT	READY	AGE
replicaset.apps/web-application-6f669d8fb6	1	1	1	14m

NAME	REFERENCE	TARGETS	MINPODS	MAXPODS	REPLICAS	AGE
horizontalpodautoscaler.autoscaling/demo-hpa	Deployment/web-application	0%/20%	1	5	1	7m59s

4.2 Let's increase the load by running the siege utility.

Terminal 1:

```
# siege -q -c 5 -t 2m http://10.105.203.79
```

Note: -q = quiet mode, -c = concurrent users (we are setting it to 5), -t = time (we are setting it to 2 minutes) and accessing one of the nodes and using the NodePort of nginx app.

Info: The controller checks the metrics every 15 seconds, as the load gradually increases, the pod will begin to autoscale-up. Continue to watch the Terminal-2 to watch the pods autoscale

4.3 Let's continue to watch the pods by executing the below command.

Terminal 2:

```
# watch kubectl get all
```

Output:

```
Every 2.0s: kubectl get all                                eoc-controller: Fri Sep  8 03:12:19 2023
```

NAME	READY	STATUS	RESTARTS	AGE
pod/web-application-6f669d8fb6-bf8n2	1/1	Running	0	17m
pod/web-application-6f669d8fb6-g96pc	1/1	Running	0	35s
pod/web-application-6f669d8fb6-h9zm5	1/1	Running	0	50s
pod/web-application-6f669d8fb6-qrdbc	1/1	Running	0	50s
pod/web-application-6f669d8fb6-rbgcr	1/1	Running	0	50s

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
service/kubernetes	ClusterIP	10.96.0.1	<none>	443/TCP	4d
service/web-service	ClusterIP	10.105.203.79	<none>	80/TCP	15m

NAME	READY	UP-TO-DATE	AVAILABLE	AGE
deployment.apps/web-application	5/5	5	5	17m

NAME	DESIRED	CURRENT	READY	AGE
replicaset.apps/web-application-6f669d8fb6	5	5	5	17m

NAME	REFERENCE	TARGETS	MINPODS	MAXPODS	REPLICAS	AGE
horizontalpodautoscaler.autoscaling/demo-hpa	Deployment/web-application	46%/20%	1	5	5	10m

Note: The siege command will complete after 2 minutes and then load will be reduce.

4.4 Let's continue to watch the pods by executing the below command.

Terminal 2:

```
# watch kubectl get all
```

Output:

```
Every 2.0s: kubectl get all                      eoc-controller: Fri Sep 8 03:16:17 2023
```

NAME	READY	STATUS	RESTARTS	AGE
pod/web-application-6f669d8fb6-bf8n2	1/1	Running	0	21m
pod/web-application-6f669d8fb6-g96pc	1/1	Running	0	4m33s
pod/web-application-6f669d8fb6-h9zm5	1/1	Running	0	4m48s
pod/web-application-6f669d8fb6-qrdbc	1/1	Running	0	4m48s
pod/web-application-6f669d8fb6-rbgcr	1/1	Running	0	4m48s

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
service/kubernetes	ClusterIP	10.96.0.1	<none>	443/TCP	4d
service/web-service	ClusterIP	10.105.203.79	<none>	80/TCP	19m

NAME	READY	UP-TO-DATE	AVAILABLE	AGE
deployment.apps/web-application	5/5	5	5	21m

NAME	DESIRED	CURRENT	READY	AGE
replicaset.apps/web-application-6f669d8fb6	5	5	5	21m

NAME	REFERENCE	TARGETS	MINPODS	MAXPODS	REPLICAS	AGE
horizontalpodautoscaler.autoscaling/demo-hpa	Deployment/web-application	0%/20%	1	5	5	14m

Note: As soon as the load is reduced, the cpu % is back to zero, but the pods will no autoscale-down immediately. Instead, it will wait for **stabilizationWindowSeconds**: which by default is 300 seconds and then scale-down.

After 5-7 minutes:

Output:

```
Every 2.0s: kubectl get all                      eoc-controller: Fri Sep 8 03:18:47 2023
```

NAME	READY	STATUS	RESTARTS	AGE
pod/web-application-6f669d8fb6-bf8n2	1/1	Running	0	23m

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
service/kubernetes	ClusterIP	10.96.0.1	<none>	443/TCP	4d
service/web-service	ClusterIP	10.105.203.79	<none>	80/TCP	22m

NAME	READY	UP-TO-DATE	AVAILABLE	AGE
deployment.apps/web-application	1/1	1	1	23m

NAME	DESIRED	CURRENT	READY	AGE
replicaset.apps/web-application-6f669d8fb6	1	1	1	23m

NAME	REFERENCE	TARGETS	MINPODS	MAXPODS	REPLICAS	AGE
horizontalpodautoscaler.autoscaling/demo-hpa	Deployment/web-application	0%/20%	1	5	1	16m

5 Cleanup.

5.1 Let's delete the hpa deployment app by executing the below commands.

```
# kubectl delete -f ~/kubernetes/hpa-deployment.yml
```

Output:

```
[root@eoc-controller ~]# kubectl delete -f ~/kubernetes/hpa-deployment.yml
deployment.apps "web-application" deleted
```

```
# kubectl delete hpa demo-hpa
```


Output:

```
[root@eoc-controller ~]#kubectl delete hpa demo-hpa
horizontalpodautoscaler.autoscaling "demo-hpa" deleted
```

5.2 Let's **delete** the **service** by executing the below commands.

```
# kubectl delete service web-service
```

Output:

```
[root@eoc-controller ~]#kubectl delete service web-service
service "web-service" deleted
```

Eyes On Cloud